

如何在MySQL中分配

innodb_buffer_pool_size

`innodb_buffer_pool_size`是整个MySQL服务器最重要的变量。

1. 为什么需要innodb_buffer_pool?

在MySQL5.5之前，广泛使用的和默认的存储引擎是MyISAM。MyISAM使用操作系统缓存来缓存数据。InnoDB需要innodb_buffer_pool中处理缓存。所以非常需要有足够的InnoDB buffer pool空间。

2. MySQL InnoDB buffer pool 里包含什么?

- 数据缓存

InnoDB数据页面

- 索引缓存

索引数据

- 缓冲数据

脏页（在内存中修改尚未刷新(写入)到磁盘的数据）

- 内部结构

如自适应哈希索引，行锁等。

3. 如何设置innodb_buffer_pool_size?

`innodb_buffer_pool_size`默认大小为128M。最大值取决于CPU的架构。在32-bit平台上，最大值为 $2^{32}-1$ ，在64-bit平台上最大值为 $2^{64}-1$ 。当缓冲池大小大于1G时，将

`innodb_buffer_pool_instances`设置大于1的值可以提高服务器的可扩展性。

大的缓冲池可以减小多次磁盘I/O访问相同的表数据。在专用数据库服务器上，可以将缓冲池大小设置为服务器物理内存的80%。

3.1 配置缓冲池大小时，请注意以下潜在问题

- 物理内存争用可能导致操作系统频繁的paging
- InnoDB为缓冲区和control structures保留了额外的内存，因此总分配空间比指定的缓冲池大小大约大10%。
- 缓冲池的地址空间必须是连续的，这在带有在特定地址加载的DLL的Windows系统上可能是一个问题。
- 初始化缓冲池的时间大致与其大小成比例。在具有大缓冲池的实例上，初始化时间可能很长。要减少初始化时间，可以在服务器关闭时保存缓冲池状态，并在服务器启动时将其还原。
 - `innodb_buffer_pool_dump_pct`: 指定每个缓冲池最近使用的页面读取和转储的百分比。范围是1到100。默认值是25。例如，如果有4个缓冲池，每个缓冲池有100个page，并且`innodb_buffer_pool_dump_pct`设置为25，则dump每个缓冲池中最近使用的25个page。

- `innodb_buffer_pool_dump_at_shutdown`: 默认启用。指定在MySQL服务器关闭时是否记录在InnoDB缓冲池中缓存的页面，以便在下次重新启动时缩短预热过程。
- `innodb_buffer_pool_load_at_startup`: 默认启用。指定在MySQL服务器启动时，InnoDB缓冲池通过加载之前保存的相同页面自动预热。通常与`innodb_buffer_pool_dump_at_shutdown`结合使用。

增大或减小缓冲池大小时，将以chunk的形式执行操作。chunk大小由

`innodb_buffer_pool_chunk_size`配置选项定义，默认值为128 MB。

缓冲池大小必须始终等于或者是`innodb_buffer_pool_chunk_size * innodb_buffer_pool_instances`的倍数。

如果将缓冲池大小更改为不等于或等于`innodb_buffer_pool_chunk_size * innodb_buffer_pool_instances`的倍数的值，

则缓冲池大小将自动调整为等于或者是`innodb_buffer_pool_chunk_size * innodb_buffer_pool_instances`的倍数的值。

`innodb_buffer_pool_size`可以动态设置，允许在不重新启动服务器的情况下调整缓冲池的大小。可以通过状态变量`Innodb_buffer_pool_resize_status`报告在线调整缓冲池大小操作的状态。

```
mysql> show status like 'Innodb_buffer_pool_resize%';
```

| Variable_name | Value |
|----------------------------------|-------|
| Innodb_buffer_pool_resize_status | |

3.2 配置示例

在以下示例中，`innodb_buffer_pool_size`设置为3G，`innodb_buffer_pool_instances`设置为8。

`innodb_buffer_pool_chunk_size`默认值为128M。

3G是有效的`innodb_buffer_pool_size`值，因为3G是`innodb_buffer_pool_instances = 8 * innodb_buffer_pool_chunk_size = 128M`的倍数

```
# mysqld --innodb_buffer_pool_size=3G --innodb_buffer_pool_instances=8 &
```

```
mysql> show variables like 'innodb_buffer_pool%';
```

| Variable_name | Value |
|--------------------------------------------------|------------------|
| <code>innodb_buffer_pool_chunk_size</code> | 134217728 |
| <code>innodb_buffer_pool_dump_at_shutdown</code> | ON |
| <code>innodb_buffer_pool_dump_now</code> | OFF |
| <code>innodb_buffer_pool_dump_pct</code> | 25 |
| <code>innodb_buffer_pool_filename</code> | ib_buffer_pool |
| <code>innodb_buffer_pool_instances</code> | 8 |
| <code>innodb_buffer_pool_load_abort</code> | OFF |

| | |
|------------------------------------|------------|
| innodb_buffer_pool_load_at_startup | ON |
| innodb_buffer_pool_load_now | OFF |
| innodb_buffer_pool_size | 3221225472 |

```
10 rows in set (0.01 sec)
```

在以下示例中，innodb_buffer_pool_size设置为3G，innodb_buffer_pool_instances设置为16。innodb_buffer_pool_chunk_size为128M。

3G不是有效的innodb_buffer_pool_size值，因为3G不是innodb_buffer_pool_instances = 16 * innodb_buffer_pool_chunk_size = 128M的倍数，可以看出innodb_buffer_pool_size的值自动调整到4GB

```
# mysqld --innodb_buffer_pool_size=3G --innodb_buffer_pool_instances=16 &
```

```
mysql> show variables like '%innodb_buffer_pool%';
```

| Variable_name | Value |
|-------------------------------------|----------------|
| innodb_buffer_pool_chunk_size | 134217728 |
| innodb_buffer_pool_dump_at_shutdown | ON |
| innodb_buffer_pool_dump_now | OFF |
| innodb_buffer_pool_dump_pct | 25 |
| innodb_buffer_pool_filename | ib_buffer_pool |
| innodb_buffer_pool_instances | 16 |
| innodb_buffer_pool_load_abort | OFF |
| innodb_buffer_pool_load_at_startup | ON |
| innodb_buffer_pool_load_now | OFF |
| innodb_buffer_pool_size | 4294967296 |

```
10 rows in set (0.01 sec)
```

3.3 在线调整InnoDB缓冲池大小

```
mysql> SET GLOBAL innodb_buffer_pool_size = 3221225472
```

3.4 监控在线缓冲池调整进度

```
mysql> SHOW STATUS WHERE Variable_name='InnoDB_buffer_pool_resize_status';
```

| Variable_name | Value |
|----------------------------------|---------------------------------------------------|
| InnoDB_buffer_pool_resize_status | Completed resizing buffer pool at 18082415:05:03. |

缓冲池大小调整进度也记录在服务器错误日志中。以下为增大时，记录的日志

2018-08-24T07:05:03.819049Z 2 [Note] InnoDB: Requested to resize buffer pool.
 (new size: 3221225472 bytes)
 2018-08-24T07:05:03.819141Z 0 [Note] InnoDB: Resizing buffer pool from
 2684354560 to 3221225472 (unit=134217728).
 2018-08-24T07:05:03.819155Z 0 [Note] InnoDB: Disabling adaptive hash index.
 2018-08-24T07:05:03.824902Z 0 [Note] InnoDB: disabled adaptive hash index.
 2018-08-24T07:05:03.824933Z 0 [Note] InnoDB: Withdrawing blocks to be shrunken.
 2018-08-24T07:05:03.824940Z 0 [Note] InnoDB: Latching whole of buffer pool.
 2018-08-24T07:05:03.824959Z 0 [Note] InnoDB: buffer pool 0 : resizing with
 chunks 5 to 6.
 2018-08-24T07:05:03.839564Z 0 [Note] InnoDB: buffer pool 0 : 1 chunks (8192
 blocks) were added.
 2018-08-24T07:05:03.839594Z 0 [Note] InnoDB: buffer pool 1 : resizing with
 chunks 5 to 6.
 2018-08-24T07:05:03.848910Z 0 [Note] InnoDB: buffer pool 1 : 1 chunks (8192
 blocks) were added.
 2018-08-24T07:05:03.849046Z 0 [Note] InnoDB: buffer pool 2 : resizing with
 chunks 5 to 6.
 2018-08-24T07:05:03.856711Z 0 [Note] InnoDB: buffer pool 2 : 1 chunks (8192
 blocks) were added.
 2018-08-24T07:05:03.856741Z 0 [Note] InnoDB: buffer pool 3 : resizing with
 chunks 5 to 6.
 2018-08-24T07:05:03.864867Z 0 [Note] InnoDB: buffer pool 3 : 1 chunks (8192
 blocks) were added.
 2018-08-24T07:05:03.864902Z 0 [Note] InnoDB: Completed to resize buffer pool
 from 2684354560 to 3221225472.
 2018-08-24T07:05:03.864915Z 0 [Note] InnoDB: Re-enabled adaptive hash index.
 2018-08-24T07:05:03.864935Z 0 [Note] InnoDB: Completed resizing buffer pool at
 180824 15:05:03.

以下为减小时，记录的日志

2018-08-24T07:10:20.666816Z 2 [Note] InnoDB: Requested to resize buffer pool.
 (new size: 2684354560 bytes)
 2018-08-24T07:10:20.666880Z 0 [Note] InnoDB: Resizing buffer pool from
 3221225472 to 2684354560 (unit=134217728).
 2018-08-24T07:10:20.666889Z 0 [Note] InnoDB: Disabling adaptive hash index.
 2018-08-24T07:10:20.673416Z 0 [Note] InnoDB: disabled adaptive hash index.
 2018-08-24T07:10:20.673508Z 0 [Note] InnoDB: Withdrawing blocks to be shrunken.
 2018-08-24T07:10:20.673519Z 0 [Note] InnoDB: buffer pool 0 : start to withdraw
 the last 8192 blocks.
 2018-08-24T07:10:20.678441Z 0 [Note] InnoDB: buffer pool 0 : withdrawing blocks.
 (8192/8192)
 2018-08-24T07:10:20.678521Z 0 [Note] InnoDB: buffer pool 0 : withdrew 8192
 blocks from free list. Tried to relocate 0 pages (8192/8192).
 2018-08-24T07:10:20.678919Z 0 [Note] InnoDB: buffer pool 0 : withdrawn target
 8192 blocks.
 2018-08-24T07:10:20.678977Z 0 [Note] InnoDB: buffer pool 1 : start to withdraw

the last **8192** blocks.

2018-08-24T07:10:20.681644Z 0 [Note] InnoDB: buffer pool **1** : withdrawing blocks.
(**8192/8192**)

2018-08-24T07:10:20.682168Z 0 [Note] InnoDB: buffer pool **1** : withdrew **8192**
blocks from free list. Tried to relocate **0** pages (**8192/8192**).

2018-08-24T07:10:20.682235Z 0 [Note] InnoDB: buffer pool **1** : withdrawn target
8192 blocks.

2018-08-24T07:10:20.682254Z 0 [Note] InnoDB: buffer pool **2** : start to withdraw
the last **8192** blocks.

2018-08-24T07:10:20.686560Z 0 [Note] InnoDB: buffer pool **2** : withdrawing blocks.
(**8192/8192**)

2018-08-24T07:10:20.686917Z 0 [Note] InnoDB: buffer pool **2** : withdrew **8192**
blocks from free list. Tried to relocate **0** pages (**8192/8192**).

2018-08-24T07:10:20.687002Z 0 [Note] InnoDB: buffer pool **2** : withdrawn target
8192 blocks.

2018-08-24T07:10:20.687010Z 0 [Note] InnoDB: buffer pool **3** : start to withdraw
the last **8192** blocks.

2018-08-24T07:10:20.690038Z 0 [Note] InnoDB: buffer pool **3** : withdrawing blocks.
(**8192/8192**)

2018-08-24T07:10:20.690373Z 0 [Note] InnoDB: buffer pool **3** : withdrew **8192**
blocks from free list. Tried to relocate **0** pages (**8192/8192**).

2018-08-24T07:10:20.690433Z 0 [Note] InnoDB: buffer pool **3** : withdrawn target
8192 blocks.

2018-08-24T07:10:20.690479Z 0 [Note] InnoDB: Latching whole of buffer pool.

2018-08-24T07:10:20.690498Z 0 [Note] InnoDB: buffer pool **0** : resizing with
chunks **6 to 5**.

2018-08-24T07:10:20.693293Z 0 [Note] InnoDB: buffer pool **0** : **1** chunks (**8192**
blocks) were freed.

2018-08-24T07:10:20.693357Z 0 [Note] InnoDB: buffer pool **1** : resizing with
chunks **6 to 5**.

2018-08-24T07:10:20.695947Z 0 [Note] InnoDB: buffer pool **1** : **1** chunks (**8192**
blocks) were freed.

2018-08-24T07:10:20.696011Z 0 [Note] InnoDB: buffer pool **2** : resizing with
chunks **6 to 5**.

2018-08-24T07:10:20.698977Z 0 [Note] InnoDB: buffer pool **2** : **1** chunks (**8192**
blocks) were freed.

2018-08-24T07:10:20.699288Z 0 [Note] InnoDB: buffer pool **3** : resizing with
chunks **6 to 5**.

2018-08-24T07:10:20.702088Z 0 [Note] InnoDB: buffer pool **3** : **1** chunks (**8192**
blocks) were freed.

2018-08-24T07:10:20.702398Z 0 [Note] InnoDB: Completed to resize buffer pool
from **3221225472 to 2684354560**.

2018-08-24T07:10:20.702413Z 0 [Note] InnoDB: Re-enabled adaptive hash index.

2018-08-24T07:10:20.703896Z 0 [Note] InnoDB: Completed resizing buffer pool at
180824 15:10:20.

4. 配置的innodb_buffer_pool_size是否合适?

当前配置的innodb_buffer_pool_size是否合适, 可以通过分析InnoDB缓冲池的性能来验证。
可以使用以下公式计算InnoDB缓冲池性能:

$Performance = innodb_buffer_pool_reads / innodb_buffer_pool_read_requests * 100$

innodb_buffer_pool_reads: 表示InnoDB缓冲池无法满足的请求数。需要从磁盘中读取。

innodb_buffer_pool_read_requests: 表示从内存中读取逻辑的请求数。

例如, 在我的服务器上, 检查当前InnoDB缓冲池的性能:

```
root@localhost [(none)] 15:35:31>show status like 'innodb_buffer_pool_read%';
```

| Variable_name | Value |
|---------------------------------------|------------|
| Innodb_buffer_pool_read_ahead_rnd | 0 |
| Innodb_buffer_pool_read_ahead | 0 |
| Innodb_buffer_pool_read_ahead_evicted | 0 |
| Innodb_buffer_pool_read_requests | 4029033624 |
| Innodb_buffer_pool_reads | 91661 |

5 rows in set (0.00 sec)

$Performance = 91661 / 4029033624 * 100 = 0.0022750120389663$

意味着InnoDB可以满足缓冲池本身的大部分请求。从磁盘完成读取的百分比非常小。因此无需增加innodb_buffer_pool_size值。

4.1 什么时候减小innodb_buffer_pool_size?

在专用MySQL服务器上, 多余的innodb_buffer内存不会有问题, 但是当使用共享服务器时, 可能会有性能影响。因为空闲内存对其他程序和操作系统很有用。

可以使用SHOW ENGINE INNODB STATUS\G命令检查内存状态:

```
mysql> show engine innodb status\G
```

...

Total large memory allocated 26386366464

Dictionary memory allocated 23826297

Buffer pool size 1572672

Free buffers 8192

Database pages 1553364

Old database pages 573246

Modified db pages 36

Pending reads 0

Pending writes: LRU 0, flush list 0, single page 0

Pages made young 881819, not young 18198964

0.02 youngs/s, 0.05 non-youngs/s

Pages read 681064, created 2749237, written 3988300

0.02 reads/s, 0.12 creates/s, 11.50 writes/s

Buffer pool hit rate 1000 / 1000, young-making rate 0 / 1000 not 0 / 1000

```
Pages read ahead 0.00/s, evicted without access 0.00/s, Random read ahead 0.00/s
LRU len: 1553364, unzip_LRU len: 0
I/O sum[5152]:cur[0], unzip sum[0]:cur[0]
...
```

Free buffers :表示有多少空闲buffer。如果 此值长时间都较高, 则可以考虑减小InnoDB缓冲池大小。

InnoDB buffer pool 命中率:

InnoDB buffer pool 命中率 = $\text{innodb_buffer_pool_read_requests} / (\text{innodb_buffer_pool_read_requests} + \text{innodb_buffer_pool_reads}) * 100$

此值低于99%, 则可以考虑增加innodb_buffer_pool_size。

5. InnoDB缓冲池状态变量有哪些?

可以运行以下命令进行查看:

```
root@localhost [(none)] 16:00:31>show global status like
'%innodb_buffer_pool_pages%';
```

| Variable_name | Value |
|----------------------------------|--------|
| Innodb_buffer_pool_pages_data | 457 |
| Innodb_buffer_pool_pages_dirty | 0 |
| Innodb_buffer_pool_pages_flushed | 36 |
| Innodb_buffer_pool_pages_free | 163363 |
| Innodb_buffer_pool_pages_misc | 0 |
| Innodb_buffer_pool_pages_total | 163820 |

6 rows in set (0.00 sec)

说明:

- **Innodb_buffer_pool_pages_data**

InnoDB缓冲池中包含数据的页数。该数字包括脏页面和干净页面。使用压缩表时, 报告的Innodb_buffer_pool_pages_data值可能大于Innodb_buffer_pool_pages_total (Bug # 59550)。

- **Innodb_buffer_pool_pages_dirty**

显示在内存中修改但尚未写入数据文件的InnoDB缓冲池数据页的数量(脏页刷新)。

- **Innodb_buffer_pool_pages_flushed**

表示从InnoDB缓冲池中刷新脏页的请求数。

- **Innodb_buffer_pool_pages_free**

显示InnoDB缓冲池中的空闲页面

- **Innodb_buffer_pool_pages_misc**

InnoDB缓冲池中的页面数量很多, 因为它们已被分配用于管理开销, 例如行锁或自适应哈希索引。此值也可以计算为Innodb_buffer_pool_pages_total - Innodb_buffer_pool_pages_free - Innodb_buffer_pool_pages_data。

- **Innodb_buffer_pool_pages_total**

InnoDB缓冲池的总大小, 以page为单位。

- **innodb_buffer_pool_reads**

表示InnoDB缓冲池无法满足的请求数。需要从磁盘中读取。

- **innodb_buffer_pool_read_requests**

它表示从内存中逻辑读取的请求数。

- **innodb_buffer_pool_wait_free**

通常，对InnoDB缓冲池的写入发生在后台。当InnoDB需要读取或创建页面并且没有可用的干净页面时，InnoDB首先刷新一些脏页并等待该操作完成。此计数器计算这些等待的实例。如果已正确设置innodb_buffer_pool_size，则此值应该很小。如果大于0，则表示InnoDB缓冲池太小。

- **innodb_buffer_pool_write_request**

表示对缓冲池执行的写入次数。

6. InnoDB缓冲池当前使用了多少实际GB内存？

通过将缓冲池中可用的数据与InnoDB页面（InnoDB缓冲池单位）大小相乘，我们可以发现InnoDB缓冲池此时正在使用的实际内存。

```
set @ibpddata = (select variable_value from information_schema.global_status
where variable_name = 'innodb_buffer_pool_pages_data');
```

```
ERROR 3167 (HY000): The 'INFORMATION_SCHEMA.GLOBAL_STATUS' feature is disabled;
see the documentation for 'show_compatibility_56'
```

#从MySQL 5.7.6开始，GLOBAL_STATUS表中提供的信息从Performance Schema获取

```
mysql> set @ibpddata = (select variable_value from
performance_schema.global_status where variable_name =
'innodb_buffer_pool_pages_data');
```

```
mysql> select @ibpddata;
```

```
+-----+
| @ibpddata |
+-----+
| 568 |
+-----+
```

```
1 row in set (0.00 sec)
```

```
mysql> set @idbpgsize = (select variable_value from
performance_schema.global_status where variable_name = 'innodb_page_size');
```

```
mysql> select @idbpgsize;
```

```
+-----+
| @idbpgsize |
+-----+
| 16384 |
+-----+
```

```
1 row in set (0.00 sec)
```



```
mysql> set @ibpsize = @ibpdata * @idbpsize / (1024*1024*1024);  
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> select @ibpsize;
```

| |
|-----------------|
| |
| -----+ |
| @ibpsize |
| -----+ |
| 0.0086669921875 |
| -----+ |

```
1 row in set (0.00 sec)
```